

Machine Learning Cheat Sheet

By Ben NG Hok Chun

Distributions

Expectation: $\mathbb{E}[x] = \int xp(x)dx$

Mean: $\mathbb{E}[\mathbf{x}]$

Variance: $\sigma^2 = \mathbb{E}[(x - \mu)^2]$

Cov(\mathbf{x}) = $\mathbb{E}[(\mathbf{x} - \mu)(\mathbf{x} - \mu)^T]$

Normal: $p(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp(-\frac{(x-\mu)^2}{2\sigma^2})$

Binomial: $p(x) = \binom{n}{k} p^k (1-p)^{n-k}$

Multinomial: $p(x) = \frac{n!}{x_1! \dots x_k!} p_1^{x_1} \dots p_k^{x_k}$

Poisson: $p(x) = \frac{\lambda^k e^{-\lambda}}{k!}$

Multivariate Normal:

$p(x) = \det(2\pi\Sigma)^{-\frac{1}{2}} \exp(-\frac{1}{2}(\mathbf{x} - \mu)^T \Sigma^{-1}(\mathbf{x} - \mu))$

Common Functions

$\varsigma(x) = \frac{1}{1+\exp(-x)}$ (Posterior class probability)

$\varsigma' = \varsigma(1 - \varsigma)$

Likelihood: $L(\theta|\mathcal{X}) = p(\mathcal{X}|\theta)$

Cross-entropy: $E(\theta|\mathcal{X}) = -\log L(\theta|\mathcal{X})$

$\sigma_i = \sigma(x_i) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$ (Posterior class probability)

$\frac{\partial \sigma_i}{\partial x_i} = \sigma_i(1 - \sigma_i)$

$\frac{\partial \sigma_i}{\partial x_j} = -\sigma_i \sigma_j, i \neq j$

Lagrangian: $\mathcal{L} = a - \sum_l \alpha_l b_l, \alpha_l \geq 0 \quad \forall l$

for minimize a subject to $b_l \geq 0 \quad \forall l$

Kernel: $K(a, b) = \phi(a)^T \phi(b)$

Polynomial kernel: $(a^T b + 1)^q$

RBF kernel: $\exp(-\frac{\|a-b\|^2}{2s^2})$ OR

$\exp(-\frac{\mathcal{D}(a,b)}{2s^2})$ for some distance function \mathcal{D}

Sigmoidal kernel: $\tanh(2a^T b + 1)$

Leaky ReLU: $f(x) = \begin{cases} x & \text{if } x \geq 0 \\ \alpha & \text{else} \end{cases}$

Exponential linear: $f(x) = \begin{cases} x & \text{if } x \geq 0 \\ \alpha(e^x - 1) & \text{else} \end{cases}$

Expectation Maximization

Expectation: $\mathcal{Q}(\Phi|\Phi^t) = \mathbb{E}[\mathcal{L}_C(\Phi|\mathcal{X}, \mathcal{Z})|\mathcal{X}, \Phi^t]$

Maximization: $\Phi^{t+1} = \arg \max_{\Phi} \mathcal{Q}(\Phi|\Phi^t)$

Deep-learning Concepts

Xavier Initialization: (for sigmoid activation):

- Zero-mean normal with variance: $\frac{2}{n_{in} + n_{out}}$

- Uniform in $[-r, r]$ where $r = \sqrt{\frac{6}{n_{in} + n_{out}}}$

He Initialization: (for ReLUs activation):

- Zero-mean normal with variance: $\frac{4}{n_{in} + n_{out}}$

- Uniform in $[-r, r]$ where $r = \sqrt{\frac{12}{n_{in} + n_{out}}}$

Batch normalization: zero-center and normalize every layer.

Dropout: probabilistically set some hidden unit to 0.

Data augmentation: Translate, scale, shift, light/dim, flip to generate new data.

AdaGrad:

$$s \leftarrow s + \nabla_w L \circ \nabla_w L$$

$$w \leftarrow w - \eta \nabla_w L \oslash \sqrt{s + \epsilon}$$

RMSProp:

$$s \leftarrow \beta s + (1 - \beta) \nabla_w L \circ \nabla_w L$$

$$w \leftarrow w - \eta \nabla_w L \oslash \sqrt{s + \epsilon}$$

Adam:

$$\Delta w \leftarrow \frac{\beta_1 \Delta w - (1 - \beta_1) \nabla_w L}{1 - \beta_1^t}$$

$$s \leftarrow \frac{\beta_2 s + (1 - \beta_2) \nabla_w L \circ \nabla_w L}{1 - \beta_2^t}$$

$$w \leftarrow w - \eta \Delta w \oslash \sqrt{s + \epsilon}$$

Derivative Rules

c	0	e^x	e^x
x	1	a^x	$\ln(a)a^x$
cx	c	$\ln(x)$	$\frac{1}{x}$
x^n	nx^{n-1}	fg	$f'g + fg'$

Integration Rules

$$\int uv dx = u \int v dx + \int u' (\int v dx) dx$$

$$\int f(g(x))g'(x) dx = \int f(u) du$$

Nonparametric Methods

Histogram Estimator:

$$\hat{p}(x) = \frac{\#\{x^{(l)} \text{ in the same bin as } x\}}{Nh}$$

$$\text{Naive Estimator: } \hat{p}(x) = \frac{\sum_l w(\frac{x-x^{(l)}}{h})}{Nh}$$

$$\text{where } w(u) = \begin{cases} 1 & \text{if } |u| < 1/2 \\ 0 & \text{else} \end{cases}$$

$$\text{Kernel Estimator: } \hat{p}(x) = \frac{\sum_l K(\frac{x-x^{(l)}}{h})}{Nh}$$

$$\text{K-nearest neighbour: } \hat{p}(x) = \frac{k}{2Nd_k(x)}$$

where $d_k(x)$ is the distance of x and k-th nearest neighbour of x

$$\text{KNN-kernel: } \hat{p}(x) = \frac{\sum_l K(\frac{x-x^{(l)}}{d_k(x)})}{Nd_k(x)}$$

Regressogram: Mean of the same bin

Running-mean smoother: Mean of the bin around x

$$\text{Kernel smoother: } \hat{g}(x) = \frac{\sum_l K(\frac{x-x^{(l)}}{h})y^{(l)}}{\sum_l K(\frac{x-x^{(l)}}{h})}$$

Running line smoother: With piecewise linear fit

Dimensionality Reduction

Feature selection: Choose k from d features

Forward search v.s. Backward search

Feature extraction: Project \mathbf{x} to \mathcal{R}^k

Principal Component Analysis:

Map \mathbf{x} to k orthogonal dimensions

$\mathbf{z}_n = \mathbf{w}_n^T \mathbf{x}$, $\text{Var}(\mathbf{z}_1) = \mathbf{w}_1^T \Sigma \mathbf{w}_1$, $\Sigma = \text{Cov}(\mathbf{x})$

Maximize $\text{Var}(\mathbf{z}_1)$ s.t. $\|\mathbf{w}_1\| = 1$ is eigenvalue of Σ

Factor Analysis: Sample \mathcal{X} , $\mathbb{E}[\mathbf{x}] = \mu$, $\text{Cov}(\mathbf{x}) = \Sigma$

Factors z_j , $\mathbb{E}[z_j] = 0$, $\text{Var}(z_j) = 1$, $\text{Cov}(\mathbf{z}) = \mathbf{I}$

Noise ϵ_i , $\mathbb{E}[\epsilon_i] = 0$, $\text{Var}(\epsilon_i) = \Psi_i$, $\text{Cov}(\epsilon) = \Psi \mathbf{I}$

$\mathbf{x} - \mu = \mathbf{V}\mathbf{z} + \epsilon$, $\Sigma = \mathbf{V}\mathbf{V}^T + \Psi$

Multidimensional Scaling:

$$\mathbf{B} = \mathbf{X}\mathbf{X}^T = \mathbf{C}\mathbf{D}\mathbf{C}^T = (\mathbf{C}\mathbf{D}^{1/2})(\mathbf{C}\mathbf{D}^{1/2})^T$$

where \mathbf{C} is eigenvectors as columns and

$\mathbf{D}^{1/2}$ is diagonal matrix of square root of eigenvalues

Drop eigenvectors with low eigenvalues in \mathbf{C} and \mathbf{D}

Linear Discriminant Analysis:

Between-class scatter: $\mathbf{w}^T \mathbf{S}_B \mathbf{w}$,

$\mathbf{S}_B = (\mathbf{m}_1 - \mathbf{m}_2)(\mathbf{m}_1 - \mathbf{m}_2)^T$, $\sum_i N_i (\mathbf{m}_i - \mathbf{m})(\mathbf{m}_i - \mathbf{m})^T$

Within-class scatter: $\mathbf{w}^T \mathbf{S}_W \mathbf{w}$,

$\mathbf{S}_W = \sum_i \mathbf{S}_i$, $\mathbf{S}_i = \sum_l y_i^{(l)} (x^{(l)} - \mathbf{m}_i)(x^{(l)} - \mathbf{m}_i)^T$

Ensemble Learning

No-free-lunch theory: No single model is the best
 Combine several simple models into group
 Voting: Convex combination of base learners
 $y = f(d_1, \dots, d_L | \Phi) = \sum_{j=1}^L w_j d_j$
 Other voting rules: Weighted sum, median, min, max, product
 Mixture of experts, Gating: $y = \sum_{j=1}^L w_j(x) d_j$
 Bayesian model combination:
 $P(C_i | x) = \sum_{\mathcal{M}_j} P(C_i | x, \mathcal{M}_j) P(\mathcal{M}_j)$,
 w_j estimates prior model probability $P(\mathcal{M}_j)$
 Bagging, Bootstrap aggregating: Base learners trained on slightly different training sets.
 Draw N from \mathcal{X} with replacement
 Boosting: Combine weak learner into strong learner
 AdaBoost: Make wrongly-labeled data have higher weight in next learner's training set

Regularization

Lasso Regression (L1 Regularization):
 • Cost: $\lambda \sum_i |w_i|$
 Ridge Regression (L2 Regularization):
 • Cost: $\lambda \sum_i w_i^2$

Problems with Machine Learning

Overfitting
 Underfitting
 Explanability
 Hardware limitation
 Time limitation
 Space/Time complexity of learning algorithm

Matrix Factorization

Given a non-negative matrix \mathbf{V} , find $\mathbf{V} \approx \mathbf{WH}$ Cost functions (Lower bound = 0 iff $\mathbf{A} = \mathbf{B}$):
 • Euclidean distance (Frobenius norm):
 $\|\mathbf{A} - \mathbf{B}\|_F^2 = \sum_{ij} (A_{ij} - B_{ij})^2$
 • Kullback-Leibler divergence:
 $D(\mathbf{A} \|\mathbf{B}) = \sum_{ij} (A_{ij} \log \frac{A_{ij}}{B_{ij}} - A_{ij} + B_{ij})$
 Optimization: convex w.r.t. to \mathbf{W} or \mathbf{H} separately, not both
 Multiplicative update rules:
 For Euclidean distance:
 $H_{a\mu} \leftarrow H_{a\mu} \frac{(\mathbf{W}^T \mathbf{V})_{a\mu}}{(\mathbf{W}^T \mathbf{WH})_{a\mu}}, W_{ia} \leftarrow W_{ia} \frac{(\mathbf{VH}^T)_{ia}}{(\mathbf{WHH}^T)_{ia}},$
 For KL-divergence:
 $H_{a\mu} \leftarrow H_{a\mu} \frac{\sum_i W_{ia} V_{i\mu} / (\mathbf{WH})_{i\mu}}{\sum_k W_{ka}},$
 $W_{ia} \leftarrow W_{ia} \frac{\sum_\mu H_{a\mu} V_{i\mu} / (\mathbf{WH})_{i\mu}}{\sum_v H_{av}},$
 Probabilistic Matrix Factorization:
 $p(R_{ij} | \mathbf{U}_i, \mathbf{V}_j, \sigma^2) = \mathcal{N}(R_{ij} | \mathbf{U}_i^T \mathbf{V}_j, \sigma^2)$
 $p(\mathbf{U}_i | \sigma_U^2) = \mathcal{N}(\mathbf{U}_i | 0, \sigma_U^2 \mathbf{I})$
 $p(\mathbf{V}_j | \sigma_V^2) = \mathcal{N}(\mathbf{V}_j | 0, \sigma_V^2 \mathbf{I})$
 MAP estimation with quadratic regularization terms:

$$E = \frac{1}{2} \sum_i \sum_j I_{ij} (R_{ij} - \mathbf{U}_i^T \mathbf{V}_j)^2 + \frac{\lambda_U}{2} \sum_i \|\mathbf{U}_i\|^2 + \frac{\lambda_V}{2} \sum_j \|\mathbf{V}_j\|^2$$

where $\lambda_X = \sigma^2 / \sigma_X^2$
 Variation: $p(R_{ij} | \mathbf{U}_i, \mathbf{V}_j, \sigma^2) = \mathcal{N}(R_{ij} | \zeta(\mathbf{U}_i^T \mathbf{V}_j), \sigma^2)$

Hidden Markov Model Definition

States: $S = \{S_1, S_2, \dots, S_N\}$
 Observation: $V = \{v_1, v_2, \dots, v_M\}$
 State transition probabilities:
 $\mathbf{A} = [a_{ij}]$ where $a_{ij} \equiv P(q_{t+1} = S_j | q_t = S_i)$
 Observation probabilities:
 $\mathbf{B} = [b_j(m)]$ where $b_j(m) \equiv P(O_t = v_m | q_t = S_j)$
 Initial state probabilities:
 $\pi = [\pi_i]$ where $\pi_i \equiv P(q_1 = S_i)$

Hidden Markov Model Algorithms

Model parameter: $\lambda = (\mathbf{A}, \mathbf{B}, \pi)$

Forward variables:

$$\alpha_t(i) \equiv P(O_1, \dots, O_t, q_t = S_i | \lambda)$$

$$\alpha_i(i) \equiv P(O_1, q_1 = S_i | \lambda) = \pi_i b_i(O_1)$$

$$\alpha_{t+1}(i) = \left[\sum_j \alpha_t(j) a_{ji} \right] b_j(O_{t+1})$$

$$P(O | \lambda) = \sum_i \alpha_T(i) \text{ time } O(N^2 T)$$

Backward variables:

$$\beta_t(i) \equiv P(O_{t+1}, \dots, O_T | q_t = S_i, \lambda)$$

$$\beta_T(i) = 1$$

$$\beta_t(i) = \sum_j a_{ij} b_j(O_{t+1}) \beta_{t+1}(j)$$

$$P(O | \lambda) = \sum_i \beta_1(i) \pi_i b_i(O_1)$$

Finding states: Viterbi Algorithm:

$$\delta_1(i) = \pi_i b_i(O_1)$$

$$\psi_i(i) = 0$$

$$\delta_t(j) = (\max_i \delta_{t-1}(i) a_{ij}) \cdot b_j(O_t)$$

$$\psi_t(j) = \arg \max_i \delta_{t-1}(i) a_{ij}$$

$$p^* = \max_i \delta_T(i)$$

$$q_T^* = \arg \max_i \delta_T(i)$$

$$q_t^* = \psi_{t+1}(q_{t+1}^*), t = T-1, \dots, 1$$

Learn model parameter: Baum-Welch algorithm:

$$\zeta_t(i, j) \equiv P(q_t = S_i, q_{t+1} = S_j | O, \lambda)$$

$$= \frac{\alpha_t(i) a_{ij} b_j(O_{t+1}) \beta_{t+1}(j)}{\sum_k \sum_l \alpha_t(k) a_{kl} b_l(O_{t+1}) \beta_{t+1}(l)}$$

$$\mathbb{E}[z_{ij}^t] = \gamma_t(i), \quad \mathbb{E}[z_{ij}^t] = \zeta_t(i, j)$$

$$\hat{a}_{ij} = \frac{\sum_{k=1}^K \sum_{t=1}^{T_k-1} \zeta_t^k(i, j)}{\sum_{k=1}^K \sum_{t=1}^{T_k-1} \gamma_t^k(i)}$$

$$\hat{b}_j(m) = \frac{\sum_{k=1}^K \sum_{t=1}^{T_k} \gamma_t^k(j) \mathbf{1}(O_t^{(k)} = v_m)}{\sum_{k=1}^K \sum_{t=1}^{T_k} \gamma_t^k(j)}$$

$$\hat{\pi}_i = \frac{\sum_{k=1}^K \gamma_1^k(i)}{K}$$